

# Kubernetes Fundamentals

## **Full Stack Bootcamp (Day 9)**

# Agenda

- Day 1
  - ES6+
- Day 2
  - React Native
- Day 3
  - Angular
- Day 4
  - Springboot
  - SpringData
- Day 5
  - JSON
  - NoSQL
- Day 6
  - Relational
- Day 7
  - Junit
  - Mockito
- Day 8
  - Docker
- Day 9
  - Kubernetes
- Day 10
  - Images and tips

# Overview

- There's no better overview than:
- <https://www.youtube.com/watch?v=4ht22ReBjno>

# Basics

- Container orchestration is the automation of much of the operational effort required to run containerized workloads and services. This includes a wide range of things software teams need to manage a container's lifecycle, including provisioning, deployment, scaling (up and down), networking, load balancing and more.
  - › Networking
  - › High availability
  - › Ease of deployment & maintenance
  - › Scalability
  - › service discovery
  - › Security & Compliance
  - › Support (Community & Enterprise)
  - › Administrative overhead

# Networking

- The container orchestrators create an internal network
- Gateways
- Exposure of ports
- Load balancers

# Microservices patterns

- Data
  - › Database per service
- Deployment
  - › Multiple services per host
- Crosscutting concerns
  - › Microservice chassis
  - › Externalized configuration
- Communication
  - › API Gateway

# Microservices patterns

- Discovery
  - › Client side discovery
  - › Server side discovery
- Circuit breaker
- Security
  - › Access token
- User Interface
  - › Server side fragment composition
  - › Client side UI composition

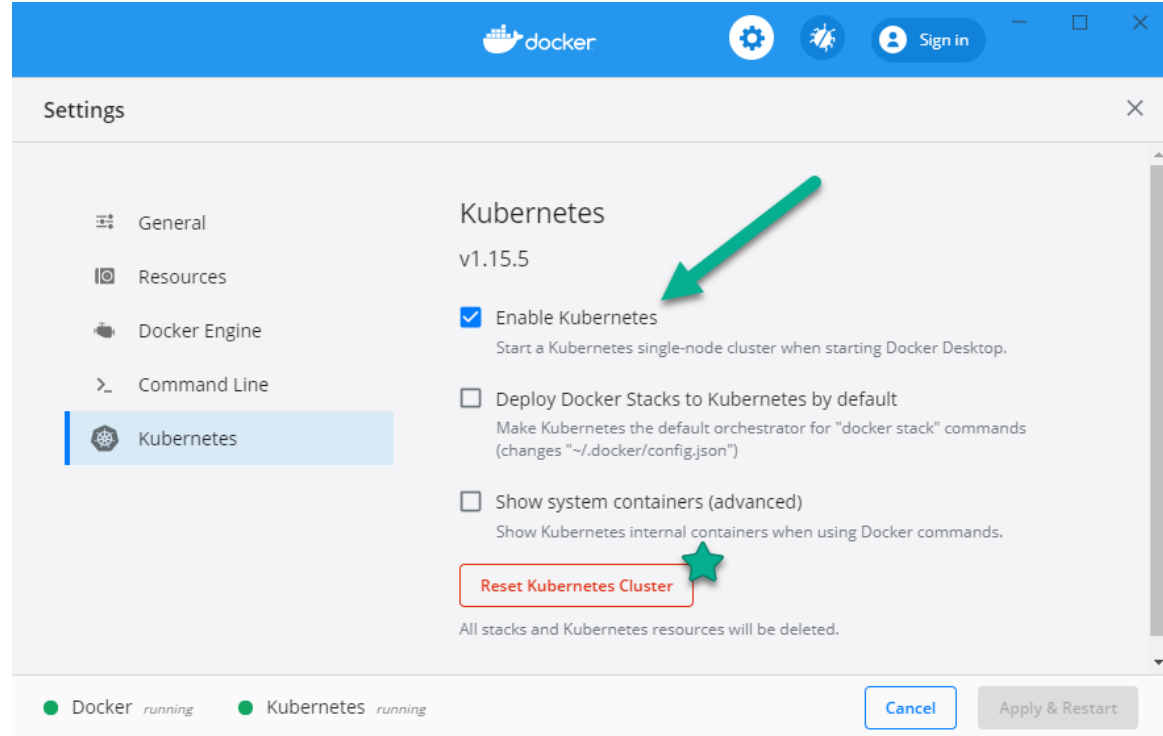
# Microservices patterns

- Observability
  - } Health check API
  - } Distributed tracing
  - } Exception tracking
  - } Application logging
- Testing
  - } Service component test
  - } Service Integration Contract Test



# Kubernetes

- Most popular orchestrator
- Provides seamlessly integration with Docker
- Basic services
  - } Pods
  - } Services
  - } Bindings
  - } Volumes
  - } Nodes
  - } Clustering



# Setup

- Find your “Kubernetes” distribution, for personal use microk8s is recommended.

<https://microk8s.io/docs/install-alternatives>

- › Download installer
- › Configure K8s
- › Test on command line

# Post installation

- On windows  
microk8s starts  
automatically
- On linux start it  
from bash  
(command line)

```
charly@worker01:~$ sudo microk8s start
Started.
charly@worker01:~$ sudo microk8s status
microk8s is running
high-availability: no
  datastore master nodes: 127.0.0.1:19001
  datastore standby nodes: none
addons:
  enabled:
    ha-cluster          # (core) Configure high availability on the current node
  disabled:
    community           # (core) The community addons repository
    dashboard           # (core) The Kubernetes dashboard
    dns                 # (core) CoreDNS
    gpu                 # (core) Automatic enablement of Nvidia CUDA
    helm                # (core) Helm 2 - the package manager for Kubernetes
    helm3               # (core) Helm 3 - Kubernetes package manager
    host-access         # (core) Allow Pods connecting to Host services smoothly
    hostpath-storage    # (core) Storage class; allocates storage from host directory
    ingress             # (core) Ingress controller for external access
    mayastor            # (core) OpenEBS MayaStor
    metallb             # (core) Loadbalancer for your Kubernetes cluster
    metrics-server      # (core) K8s Metrics Server for API access to service metrics
    prometheus          # (core) Prometheus operator for monitoring and logging
    rbac                # (core) Role-Based Access Control for authorisation
    registry            # (core) Private image registry exposed on localhost:32000
    storage              # (core) Alias to hostpath-storage add-on, deprecated
```

# Enabling dashboard

- On Linux super user powers are required

```
charly@worker01:~$ sudo microk8s enable dashboard
Infer repository core for addon dashboard
Enabling Kubernetes Dashboard
Infer repository core for addon metrics-server
Enabling Metrics-Server
serviceaccount/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
service/metrics-server created
deployment.apps/metrics-server created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
clusterrolebinding.rbac.authorization.k8s.io/microk8s-admin created
Metrics-Server is enabled
Applying manifest
serviceaccount/kubernetes-dashboard created
service/kubernetes-dashboard created
secret/kubernetes-dashboard-certs created
secret/kubernetes-dashboard-csrf created
secret/kubernetes-dashboard-key-holder created
configmap/kubernetes-dashboard-settings created
role.rbac.authorization.k8s.io/kubernetes-dashboard created
clusterrole.rbac.authorization.k8s.io/kubernetes-dashboard created
rolebinding.rbac.authorization.k8s.io/kubernetes-dashboard created
clusterrolebinding.rbac.authorization.k8s.io/kubernetes-dashboard created
deployment.apps/kubernetes-dashboard created
service/dashboard-metrics-scraper created
deployment.apps/dashboard-metrics-scraper created

If RBAC is not enabled access the dashboard using the default token retrieved with:

token=$(microk8s kubectl -n kube-system get secret | grep default-token | cut -d " " -f1)
microk8s kubectl -n kube-system describe secret $token

In an RBAC enabled setup (microk8s enable RBAC) you need to create a user with restricted
permissions as shown in:
https://github.com/kubernetes/dashboard/blob/master/docs/user/access-control/creating-sample-user.md
```

# Access Token

```
charly@worker01:~$ sudo microk8s dashboard-proxy
Checking if Dashboard is running.
Infer repository core for addon dashboard
Waiting for Dashboard to come up.
Create token for accessing the dashboard
secret/microk8s-dashboard-token created
Waiting for secret token (attempt 0)
Dashboard will be available at https://127.0.0.1:10443
Use the following token to login:
eyJhbGciOiJIUzI1NiIsImtpZCI6InZmVXBYWVwcm9udGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXJvbmVudC9uYWw1c3BhY2UiOiJrdWJlLnN5c3RlbSI9Imt1YmVybmV0ZXMuaW8vc2VydmljZWZjY291bnQvc2VudC5yZWw1IjoizGVmYXVsdCIsImt1YmVybmV0ZXMuaW8vc2VydmljZWZjY291bnQvc2VudC5yZWw1IiwidmVudG9rZW4iLCJrdWJlLnNlcnZpY2VhY2NvdW50L3NlcnZpY2UtYWNjb3VudC5yZWw1IjoizGVmYXVsdCIsImt1YmVybmV0ZXMuaW8vc2VydmljZWZjY291bnQvc2VudC5yZWw1IiwidmVudG9rZW4iLCJrdWJlLnNlcnZpY2UtYWNjb3VudDprdwJlLnN5c3RlbTpkZWZhdWx0In0.Ns1FJder2TtmlStA0HKSGdDBwc1nM227vCZ7gQPSCPss17LGCHU5cCqQ0IvihNqobS110jUtx3AQojc35TDdKShn5B_Obd6aYPGIWKhsXglSQCYS13pcjGR16RVOb0ewKmLHuwHgPQQ2BzRBP6QRLC0PBu5kz2JK9-w-RX6ZZLVzEg87IN2iPozq1J6SyO49PEswi_redW_SKRh9_FcIUZPeE8v_cmjF8yKBGNm4llvvN2MVMPifDLzVZm0QtG6Vo41RAIVXMEis1-zCtrlRpQTNO0cWE_cv-YwPjhROe5E2ejNvb_P6HMqQ0HfPYFCOdLgx6LK0xIA7xNMS6fpxIQ
```

# Autosigned dashboard



## Advertencia: Riesgo potencial de seguridad a continuación

Firefox ha detectado una potencial amenaza y no ha continuado a 10.152.183.64. Si visitas este sitio, los atacantes podrían intentar robar tu información como tus contraseñas, correo o datos de tu tarjeta de crédito.

[Saber más...](#)

**Volver (recomendado)**

Avanzado...

10.152.183.64 usa un certificado de seguridad no válido.

El certificado no es de confianza porque está autofirmado.

Código de error: [MOZILLA\\_PKIX\\_ERROR\\_SELF\\_SIGNED\\_CERT](#)

[Ver certificado](#)

**Volver (recomendado)**

Aceptar el riesgo y continuar

# Dashboard access

- **sudo microk8s dashboard-proxy**
- Copy token

## Kubernetes Dashboard

☒ Token

Every Service Account has a Secret with valid Bearer Token that can be used to log in to Dashboard. To find out more about how to configure and use Bearer Tokens, please refer to the [Authentication](#) section.

☐ Kubeconfig

Please select the kubeconfig file that you have created to configure access to the cluster. To find out more about how to configure and use kubeconfig file, please refer to the [Configure Access to Multiple Clusters](#) section.

Enter token \*

Sign in

# Dashboard

The image shows a screenshot of the Kubernetes Dashboard interface. The top navigation bar includes the Kubernetes logo, a namespace selector dropdown set to 'default', a search bar, and icons for adding resources, notifications, and user profile. The left sidebar contains a menu with categories: Workloads (Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets), Service (Ingresses, Services), and Config and Storage. The main content area is titled 'Pods' and displays the message 'There is nothing to display here' and 'No resources found.'.

Annotations with arrows point to the following elements:

- Namespace selector**: Points to the 'default' dropdown menu.
- Create resource**: Points to the '+' icon in the top right corner.
- Monitoring and working zone**: Points to the main content area.
- Activities Menu**: Points to the left sidebar menu.



# Pods



kube-system

Search



Workloads > Pods

Workloads <sup>N</sup>

Cron Jobs

Daemon Sets

Deployments

Jobs

**Pods**

Replica Sets

Replication Controllers

Stateful Sets

Service <sup>N</sup>

Ingresses

Services

Config and Storage

Config Maps <sup>N</sup>


Persistent Volume Claims <sup>N</sup>

Secrets <sup>N</sup>

CPU Usage



Memory Usage



Pods

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created ↑
<a href="#">kubernetes-dashboard-765646474b-ws6s2</a>	kubernetesui/dashboard:v2.3.0	k8s-app: kubernetes-dashb oard pod-template-hash: 76564 6474b	worker01	Running	0	8.00m	27.64Mi	7 hours ago
<a href="#">dashboard-metrics-scraper-6b6f796c8d-66s9d</a>	kubernetesui/metrics-scraper:v1.0.6	k8s-app: dashboard-metric s-scraper pod-template-hash: 6b6f79 6c8d	worker01	Running	0	3.00m	8.68Mi	7 hours ago



# Pod details

The screenshot displays the Kubernetes dashboard interface. The top navigation bar shows 'Workloads > Pods > kubernetes-dashboard-765646474b-ws6s2'. The left sidebar contains a 'CPU Usage' graph and a 'Metadata' section. The main content area is divided into two panels: 'Resource information' and 'Conditions'.

**Resource information**

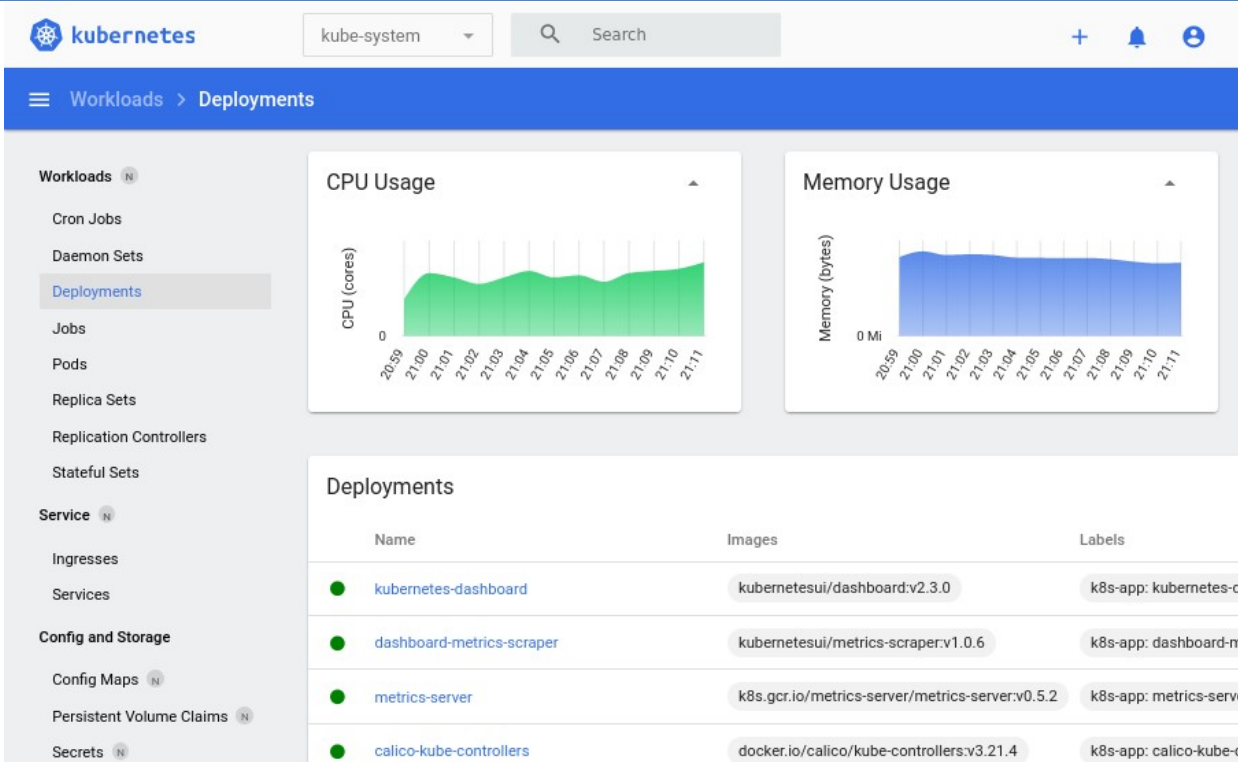
Node	Status	IP	QoS Class	Restarts	Service Account
worker01	Running	10.1.5.3	BestEffort	0	kubernetes-dashboard

**Conditions**

Type	Status	Last probe time	Last transition time	Reason
Initialized	True	-	7 hours ago	-
Ready	True	-	7 hours ago	-
ContainersReady	True	-	7 hours ago	-
PodScheduled	True	-	7 hours ago	-

- Monitoring CPU and memory usage
- Metadata
- Resource information
- Conditions
- Controlled by
- Persistence volumes claims
- Events
- Containers
  - Image
  - Status
  - Arguments
  - Mounts
  - Security Context
  - Liveness probe

# Deployments



- pods. In short, a pod is the core building block for running applications in a Kubernetes cluster; a deployment is a management tool used to control the way pods behave.

# Hello miniKube

kube-system

Q

Search

+

🔔

👤

charly@worker01:~\$ sudo microk8s kubectl get services --namespace kube-system

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
metrics-server	ClusterIP	10.152.183.181	<none>	443/TCP	9h
kubernetes-dashboard	ClusterIP	10.152.183.64	<none>	443/TCP	9h
dashboard-metrics-scraper	ClusterIP	10.152.183.56	<none>	8000/TCP	9h
hello-minikube	LoadBalancer	10.152.183.30	<pending>	3030:30911/TCP	18m

Create from input

Create from file

Create from form

Enter VAML or JSON content specifies the resources to create to the currently selected namespace. [Learn more](#)

kube-public

Q

Search

+

🔔

👤

1

Create from input

Create from file

Create from form

App name \*

hello-minikube

14 / 24

Container image \*

docker/getting-started

Number of pods \*

3

Service \*

External

Port \*

3030

Target port \*

80

Protocol \*

TCP

Port

Target port

Protocol \*

TCP

Deploy

Cancel

Show advanced options

An 'app' label with this value will be added to the Deployment and Service that get deployed. [Learn more](#)

Enter the URL of a public image on any registry, or a private image hosted on Docker Hub or Google Container Registry. [Learn more](#)

A Deployment will be created to maintain the desired number of pods across your cluster. [Learn more](#)


Optionally, an internal or external Service can be defined to map an incoming Port to a target Port seen by the container. [Learn more](#)

# Hello Minikube

The screenshot shows a web browser window with the address bar displaying `10.152.183.30:3030/tutorial/`. The browser tabs include 'Most Visited', 'Bienvenidos a Beisból...', 'Virtual University', 'Packages Search - pk...', and 'Welcom'. The main content area shows the 'docker Labs Getting Started' page. Below this, the Kubernetes dashboard is visible, showing the 'kube-system' namespace and a search bar. The 'Services' page is selected, displaying a table of services.

Name	Labels	Type	Cluster IP	Internal Endpoints
hello-minikube	k8s-app: hello-minikube	LoadBalancer	10.152.183.30	hello-minikube.kube-system:3030 TCP hello-minikube.kube-system:30911 TCP

# Configuration



- Service
- Pod
- DNS
- Deployments
- Namespaces
- Nodes

# Clusters

- A Kubernetes cluster is a set of nodes that run containerized applications.
- Containerizing applications packages an app with its dependences and some necessary services. They are more lightweight and flexible than virtual machines.

# Clusters

- A Kubernetes cluster contains six main components:
  - **API server:** Exposes a REST interface to all Kubernetes resources. Serves as the front end of the Kubernetes control plane.
  - **Scheduler:** Places containers according to resource requirements and metrics. Makes note of Pods with no assigned node, and selects nodes for them to run on.
  - **Controller manager:** Runs controller processes and reconciles the cluster's actual state with its desired specifications. Manages controllers such as node controllers, endpoints controllers and replication controllers.
  - **Kubelet:** Ensures that containers are running in a Pod by interacting with the Docker engine , the default program for creating and managing containers. Takes a set of provided **PodSpecs** and ensures that their corresponding containers are fully operational.
  - **Kube-proxy:** Manages network connectivity and maintains network rules across nodes. Implements the Kubernetes Service concept across every node in a given cluster.
  - **Etcd:** Stores all cluster data. Consistent and highly available Kubernetes backing store.



# Clusters

- Automation occurs via the Pod Lifecycle Event Generator, or PLEG. These automatic tasks can include:
  - Starting and restarting containers
  - Adjusting the number of replicas for an application
  - Validating container images
  - Launching and managing containers
  - Implementing updates and rollbacks

# Stateless Apps

- No storage assigned
- The app state is lost when **pod**, **namespace** or **kubernetes instance** is restarted
- Exposed to “the world”

# Statefull Apps

- Creating a StatefulSet
- Begin by creating a StatefulSet using the example below. It is similar to the example presented in the StatefulSets concept. It creates a headless Service, nginx, to publish the IP addresses of Pods in the StatefulSet, web.

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  ports:
    - port: 80
      name: web
  clusterIP: None
  selector:
    app: nginx
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: "nginx"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
```

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
      - name: nginx
        image: k8s.gcr.io/nginx-slim:0.8
        ports:
          - containerPort: 80
            name: web
        volumeMounts:
          - name: www
            mountPath: /usr/share/nginx/html
    volumeClaimTemplates:
      - metadata:
          name: www
        spec:
          accessModes: [ "ReadWriteOnce" ]
          resources:
            requests:
              storage: 1Gi
```

# State-full Apps

- Using the yaml file create the state-full app

```
kubectl apply -f web.yaml
```

- Review service status

```
kubectl get service nginx
```

- StatefulSets concept, the Pods in a StatefulSet have a sticky, unique identity. This identity is based on a unique ordinal index that is assigned to each Pod by the StatefulSet controller.

```
kubectl get pods -l app=nginx
```

- Each Pod has a stable hostname based on its ordinal index. Use **kubectl** exec to execute the **hostname** command in each Pod:

```
for i in 0 1; do kubectl exec "web-$i" -- sh -c 'hostname'; done
```

- Get the PersistentVolumeClaims for web-0 and web-1:

```
kubectl get pvc -l app=nginx
```

# Statefull Apps

- The **StatefulSet** controller created two **PersistentVolumeClaims** that are bound to two **PersistentVolumes**.
- As the cluster used in this tutorial is configured to dynamically provision **PersistentVolumes**, the **PersistentVolumes** were created and bound automatically.
- The **NGINX** webserver, by default, serves an index file from `/usr/share/nginx/html/index.html`. The **volumeMounts** field in the **StatefulSet's** spec ensures that the `/usr/share/nginx/html` directory is backed by a **PersistentVolume**.
- Write the Pods' hostnames to their `index.html` files and verify that the **NGINX** webserver serves the hostnames:

```
for i in 0 1; do kubectl exec "web-$i" -- sh -c 'echo "${hostname}" > /usr/share/nginx/html/index.html'; done
```

```
for i in 0 1; do kubectl exec -i -t "web-$i" -- curl http://localhost/; done
```

# Stateful Apps

- Scaling a **StatefulSet** refers to increasing or decreasing the number of replicas. This is accomplished by updating the replicas field. You can use either **kubectl** scale or **kubectl** patch to scale a **StatefulSet**.
- Scaling Up
  - In one terminal window, watch the Pods in the **StatefulSet**:  
**kubectl get pods -w -l app=nginx**
  - In another terminal window, use **kubectl** scale to scale the number of replicas to 5:
    - **kubectl scale sts web --replicas=5**

# Services

Kubernetes networking addresses four concerns:

- Containers within a Pod use networking to communicate via loopback.
- Cluster networking provides communication between different Pods.
- The Service resource lets you expose an application running in Pods to be reachable from outside your cluster.
- You can also use Services to publish services only for consumption inside your cluster.

# Services

- A service is an abstract way to expose an application running on a set of Pods as a network service.
- With Kubernetes you don't need to modify your application to use an unfamiliar service discovery mechanism. Kubernetes gives Pods their own IP addresses and a single DNS name for a set of Pods, and can load-balance across them
- If you're able to use Kubernetes APIs for service discovery in your application, you can query the API server for Endpoints, that get updated whenever the set of Pods in a Service changes.
- For non-native applications, Kubernetes offers ways to place a network port or load balancer in between your application and the backend Pods.



# Storage

- Volumes
  - On-disk files in a container are ephemeral, which presents some problems for non-trivial applications when running in containers. One problem is the loss of files when a container crashes. The kubelet restarts the container but with a clean state. A second problem occurs when sharing files between containers running together in a Pod. The Kubernetes volume abstraction solves both of these problems
- Type of volumes
  - Persistent
  - Projected
  - Ephemeral

# Volumes

- A PersistentVolume (PV) is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using Storage Classes. It is a resource in the cluster just like a node is a cluster resource. PVs are volume plugins like Volumes, but have a lifecycle independent of any individual Pod that uses the PV. This API object captures the details of the implementation of the storage, be that NFS, iSCSI, or a cloud-provider-specific storage system.
- A PersistentVolumeClaim (PVC) is a request for storage by a user. It is similar to a Pod. Pods consume node resources and PVCs consume PV resources. Pods can request specific levels of resources (CPU and Memory). Claims can request specific size and access modes (e.g., they can be mounted ReadWriteOnce, ReadOnlyMany or ReadWriteMany, see AccessModes).

# Volumes

- A projected volume maps several existing volume sources into the same directory.
- Currently, the following types of volume sources can be projected:
  - secret
  - downwardAPI
  - configMap
  - serviceAccountToken
- All sources are required to be in the same namespace as the Pod.

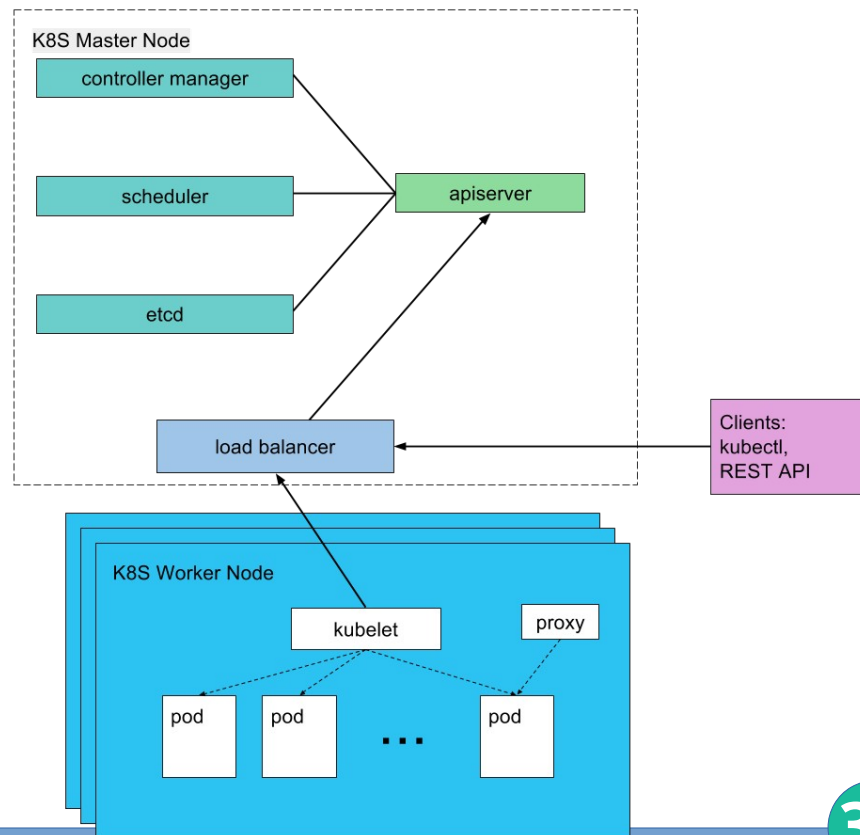
# Volumes

- Some application need additional storage but don't care whether that data is stored persistently across restarts. For example, caching services are often limited by memory size and can move infrequently used data into storage that is slower than memory with little impact on overall performance.
- Other applications expect some read-only input data to be present in files, like configuration data or secret keys.
- Ephemeral volumes are designed for these use cases.

# High availability

- Kubernetes High Availability ensures that Kubernetes and its supporting components have no single point of failure. A single master cluster is vulnerable to failure, but a multi-master cluster uses many master nodes, each having access to the same worker nodes.
- Create Highly Available Kubernetes Clusters
  - Make the master node services reliable.
  - Set up a redundant storage layer for etcd.
  - Use a highly available load balancer for the Kubernetes API services.
  - Setup multiple master nodes and configure a master election strategy.

Kubernetes Cluster



# Day 9 summary

- Kubernetes overview
- Microservices patterns
- Setup Configuration
- Dashboard
- Image usage
- Clusters
- Stateless and Statefull apps
- Services
- Storage
- High availability